

Imran Nazar

# The smallest .nds file

Articles

Programming

Fiction

Get in touch

Have you ever wondered what exactly is inside a Nintendo DS ROM file, and why the simple DS demos are so much larger than their GBA equivalents? Some people have, and this page documents their exploits.

## Introduction

It started innocently enough. I was looking for a small ROM which would be used to test the framebuffer display mode of DSEmu, a Nintendo DS emulator. LiraNuna agreed to put a small C demo together, to fill the 'main' screen with red, demonstrating the framebuffer's use. When compiled and spliced up, the .nds ended up at around 7.5KB.

That, LiraNuna thought, was a bit large for something that did so little as his demo evidently did. Stepping through with DSEmu's debugger, I noticed a whole lot of code being run which wasn't strictly required:



setting up cache parameters and the stack, clearing out regions of memory, and such like. Referred to as the `crto`, this code is inserted into every project, to safeguard the execution environment.

Furthermore, there was the standard ARM7 code also inserted into the `.nds` file, which does such things as set up the touchscreen. All this, we thought, was a bit over-the-top for a demo that was literally doing almost nothing. So, the cut-down began.

### **Early stages: Chopping the ARM7**

First off, LiraNuna thought the functionality of the ARM7 wasn't particularly required for this demo. So, the thought process went, why not simply tell that 'sub' CPU to enter an infinite loop and not do anything? The reasoning was sound, and so the ARM7 source file was replaced with a simple assembly file, looking something like this.

#### ***ARM7 cut-down: Infinite loop***

```
main: b main
```

Once put together, that reduced the size of the overall `.nds` file by quite a way; down to approximately 5KB. However, I still thought that was a touch large. A quick peek into the `.nds` file showed why that was: the sub CPU, just like the main CPU, has a `crto` automatically inserted by the build process, and this made up the vast majority of the ARM7 portion of the `.nds` file.

Therefore, LiraNuna took the step of subverting a part of the build process, by deleting the result of the ARM7 compilation, and replacing it with a straight binary file, encoding the infinite-loop opcode.

#### ***ARM7 cut-down: Final binary***

```
0000    FE FF FF EA
```



That left the overall binary at around 4KB. Still plenty of room for improvement, I thought.

### **The Next Step: Assembly**

The main code was still in C, and compiled to Thumb binary. Stepping through that in DSemu's debugger, I noticed a few odd things introduced by the compiler, that seemed to do very little; values being left-shifted and then right-shifted again, to no overall effect, and similar oddities. So, the next logical step was to write that portion without the intervention of the compiler, in assembly.

LiraNuna put together a first attempt at an assembly version of the program, as follows.

#### ***ARM9 cut-down: First run***

```
main:
    @ sets POWER_CR
    mov r0, #0x4000000
    orr r0, r0, #0x300
    orr r0, r0, #0x4
    mov r1, #0x3
    str r1, [r0]

    @ sets mode
    mov r0, #0x04000000
    mov r1, #0x00020000
    str r1, [r0]

    @ sets VRAM bank a
```



```
mov r0, #0x04000000
add r0, r0, #0x240
mov r1, #0x80
strb r1, [r0]

        @ loop
mov r0, #0x06800000
mov r1, #0x1F
orr r1, r1, #0x8000
mov r2, #0x18000

filloop:
    strh r1, [r0], #0x1
    subs r2, r2, #0x1
    bne filloop

lforever:
b lforever
```

When compiled up, that definitely made a difference; the overall ROM size dropped to approximately 1.5KB. However, I started to have an inkling that we could do better. And that's when pepsiman piped up with a suggestion: place the code *inside* the .nds header.

### **Going deeper: Inside the .nds file**

What did pepsiman mean by that? In order to understand that, it's important to know what a .nds ROM looks like, on the inside.



File offset	Component
0000	NDS ROM header (512 bytes)
0200	ARM9 binary
0200+ARM9	ARM7 binary
0200+both	Optional file table

The conventional layout dictates that the main CPU's binary be placed after the header, and the sub CPU's binary after that. However, that doesn't have to hold true all the time; the order can be swapped, blank space can be inserted between the binaries, or after them.

That's all well and good, but *inside* the header? In order to understand that, it's required to look inside that top chunk of the file: the ROM header.

**Header structure: ndstool sample output**

```

0x00  Game title
0x0C  Game code          #####
0x10  Maker code
0x12  Unit code          0x00
0x13  Device type        0x00
0x14  Device capacity    0x00 (1 Mbit)
0x15  (8 bytes blank space)
0x1E  ROM version        0x00
0x1F  reserved           0x04
0x20  ARM9 ROM offset    0x200
0x24  ARM9 entry address 0x2000000
0x28  ARM9 RAM address   0x2000000

```



0x2C	ARM9 code size	0x3A0
0x30	ARM7 ROM offset	0x600
0x34	ARM7 entry address	0x3800000
0x38	ARM7 RAM address	0x3800000
0x3C	ARM7 code size	0x8
0x40	File name table offset	0x608
0x44	File name table size	0x9
0x48	FAT offset	0x614
0x4C	FAT size	0x0
0x50	ARM9 overlay offset	0x0
0x54	ARM9 overlay size	0x0
0x58	ARM7 overlay offset	0x0
0x5C	ARM7 overlay size	0x0
0x60	ROM control info 1	0x00586000
0x64	ROM control info 2	0x001808F8
0x68	Icon/title offset	0x0
0x6C	Secure area CRC	0x0000 (-, homebrew)
0x6E	ROM control info 3	0x0000
0x70	(16 bytes blank space)	
0x80	Application end offset	0x00000000
0x84	ROM header size	0x00000200
0x88	(36 bytes blank space)	
0xAC	PassMe autoboot detect	0x53534150 ("PASS")
0xB0	(16 bytes blank space)	
0xC0	Nintendo Logo (156 bytes)	
0x15C	Logo CRC	0x9E1A (OK)
0x15E	Header CRC	0xC9D3 (OK)
0x160	(160 bytes blank space)	



The entries highlighted red indicate regions of empty space in the header structure. These are normally left behind during the construction of the format, to allow for expansion. In this case, however, it's possible to make use of the blank regions in the header for the purposes of holding code.

From looking at the above output, it's simple to see that the structure of the .nds file as a whole is dictated by the entries in this header. The fact that the ARM9 binary follows the header is simply due to the setting of "ARM9 ROM offset" to 0x200, which is the first byte in the file after the header. Similarly, the ARM7 code following the ARM9 is a simple effect of the "ARM7 ROM offset" being set to 0x600, which corresponds to an offset in the file of 1.5KB.

Simply by changing the "ROM offset" values in this header, it's possible to change the point from which the code for the CPUs is loaded, from the default location after the header to somewhere inside the header; overwrite the zeros in that position with ARM opcodes, and load from there. It seemed a good idea by pepsiman, and viable.

LiraNuna's ARM9 code seemed quite short, but I thought I could go one better, shrinking the code down further.

#### ***ARM9 cut-down: Second run***

```
main:
    mov r0,#0x04000000      ; I/O space offset
    mov r1,#0x3             ; Both screens on
    mov r2,#0x00020000      ; Framebuffer mode
    mov r3,#0x80            ; VRAM bank A enabled, LCD

    str r1,[r0, #0x304]      ; Set POWERCNT
    str r2,[r0]              ;     DISPCNT
    str r3,[r0, #0x240]      ;     VRAMCNT_A
```



```

    mov r0,#0x06800000      ; VRAM offset
    mov r1,#31              ; Writing red pixels
    mov r2,#0xC000          ; 96k of them

lp: strh r1,[r0],#2         ; Write a pixel
    subs r2,r2,#1           ; Move along one
    bne lp                  ; And loop back if not done

nf: b nf                    ; Sit in an infinite loop to finish

```

Once assembled, this code ended up looking like the following.

***ARM9 cut-down: Assembled binary***

0000	01 03 A0 E3	03 10 A0 E3	02 28 A0 E3	80 30 A0 E3
0010	04 13 80 E5	00 20 80 E5	40 32 80 E5	1A 05 A0 E3
0020	1F 10 A0 E3	03 29 A0 E3	B2 10 C0 E0	01 20 52 E2
0030	FC FF FF 1A	FE FF FF EA		

Definitely a little smaller; now the matter remained of where to put it, along with the ARM7 binary of one opcode (EAfffffE). The ARM7 was simple enough: the first region of blank space, 8 bytes, was ample space to place this opcode. The ARM7 offset was changed, the size changed to 4, and that part was done.

The ARM9 code was similarly simple to place in: the 160 bytes of free space at the end of the header seemed more than enough to stash the binary, and all that remained was to modify the ARM9 ROM offset and size.



And that, it seemed, was that. All the code fit comfortably into the header, and the final .nds was just 512 bytes in size. Surely that was all that could be done? Not quite.

### **To the core: Repositioning**

As it turns out, not all 512 bytes of the header are used. The 160 bytes on the end are in the header simply by convention; one might as well say that the .nds file consists of a 352-byte header, 160 bytes of padding, and then the two CPU binaries. Was it possible to fit the 56-byte ARM9 binary somewhere else inside the header, and eliminate this padding?

I started by changing the "header size" field at 0x84 to reflect the new size of the header, which would be 0x160 bytes. Then, I started inserting the opcodes, until I had something like this.

#### ***ARM9 placement: Within the header***

0070	01 03 A0 E3	03 10 A0 E3	02 28 A0 E3	80 30 A0 E3
0080	00 00 00 00	A0 01 00 00	04 13 80 E5	00 20 80 E5
0090	40 32 80 E5	1A 05 A0 E3	1F 10 A0 E3	03 29 A0 E3
00A0	B2 10 C0 E0	01 20 52 E2	FC FF FF 1A	50 41 53 53
00B0	FE FF FF EA	00 00 00 00	00 00 00 00	00 00 00 00

The fields in the header at 0x80, 0x84 and 0xAC can be seen, nestled within the ARM9 code. Now, this is quite a problem; if those values correspond to valid opcodes, they may be executed, and that might prove disastrous for the state of the program.

A disassembly was called for. I loaded up the new binary in DSemu, and the debugger gave the following output:



### **ARM9 cut-down: Code after insertion**

```
mov r0, #0x04000000
mov r1, #0x3
mov r2, #0x00020000
mov r3, #0x80
andeq r0, r0, r0
andeq r0, r0, r0, lsr #3
str r1, [r0, #0x304]
str r2, [r0]
str r3, [r0, #0x240]
mov r0, #0x06800000
mov r1, #31
mov r2, #0xC000
lp: strh r1, [r0], #2
    subs r2, r2, #1
    bne lp
    cmppls r3, #0x14
nf: b nf
```

It seems I was fortunate. The first two AND statements will never be executed, since they depend on the ZERO flag being set, and said flag is not set by the instructions above. As for the CMP, it slots into place after the VRAM-writing loop, which is indeed fortunate; if the CMP had fallen before the BNE, the loop may have executed forever, eventually running out of VRAM to write to.

Surprisingly fortunate, I thought; I hadn't planned for such a consequence, and it had simply come about due to the size and structure of the code. Either way, I wasn't about to complain.



## Conclusion

So, there we have it. The smallest .nds file you're ever likely to see, which still does something. The ARM7 sticks itself into an infinite loop, and the ARM9 fills the main-core framebuffer with red before entering its own infinite loop. I eventually got my wish, of a small framebuffer-testing demo, but it was fun to get there.

### *Final binary: 352 bytes*

0000	4E 44 53 2E	54 69 6E 79	46 42 00 00	23 23 23 23	NDS.TinyFB...###
0010	00 00 00 00	00 00 FE FF	FF EA 00 00	00 00 00 04	.....
0020	70 00 00 00	00 00 00 02	00 00 00 02	44 00 00 00	p.....D...
0030	16 00 00 00	00 00 80 03	00 00 80 03	04 00 00 00	.....
0040	A0 01 00 00	00 00 00 00	A0 01 00 00	00 00 00 00	.....
0050	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00	.....
0060	00 60 58 00	F8 08 18 00	00 00 00 00	00 00 00 00	.`X.....
0070	01 03 A0 E3	03 10 A0 E3	02 28 A0 E3	80 30 A0 E3	..... (...0..
0080	00 00 00 00	A0 01 00 00	04 13 80 E5	00 20 80 E5	..... ..
0090	40 32 80 E5	1A 05 A0 E3	1F 10 A0 E3	03 29 A0 E3	@2.....) ..
00A0	B2 10 C0 E0	01 20 52 E2	FC FF FF 1A	50 41 53 53	..... R.....PASS
00B0	FE FF FF EA	00 00 00 00	00 00 00 00	00 00 00 00	.....
00C0	C8 60 4F E2	01 70 8F E2	17 FF 2F E1	12 4F 11 48	.`O..p..../.O.H
00D0	12 4C 20 60	64 60 7C 62	30 1C 39 1C	10 4A 00 F0	.L `d` b0.9..J..
00E0	14 F8 30 6A	80 19 B1 6A	F2 6A 00 F0	0B F8 30 6B	..0j...j.j....0k
00F0	80 19 B1 6B	F2 6B 00 F0	08 F8 70 6A	77 6B 07 4C	...k.k....pjwk.L
0100	60 60 38 47	07 4B D2 18	9A 43 07 4B	92 08 D2 18	` `8G.K...C.K....
0110	0C DF F7 46	04 F0 1F E5	00 FE 7F 02	F0 FF 7F 02	...F.....



0120	F0 01 00 00	FF 01 00 00	00 00 00 04	00 00 00 00	.....
0130	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00	.....
0140	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00	.....
0150	00 00 00 00	00 00 00 00	00 00 00 00	1A 9E 7B EB	.....{.

<http://imrannazar.com/content/files/TinyFB.nds>

Two9A, with thanks to LiraNuna and pepsiman

*Article dated: 22nd Sep 2006*

Content copyright Imran Nazar, 2005-2014.  
Design and imagery copyright Imran Nazar, 2008-2011; "Parchment" used by license from [sxc](#).